Separate configured-cache cohorts — c1, 256 B and 8 KiB
Matched repeats r0, r1

TH TCP · TH EFA · TH EFA direct · TH EFA native

Coordinator cache: API p50

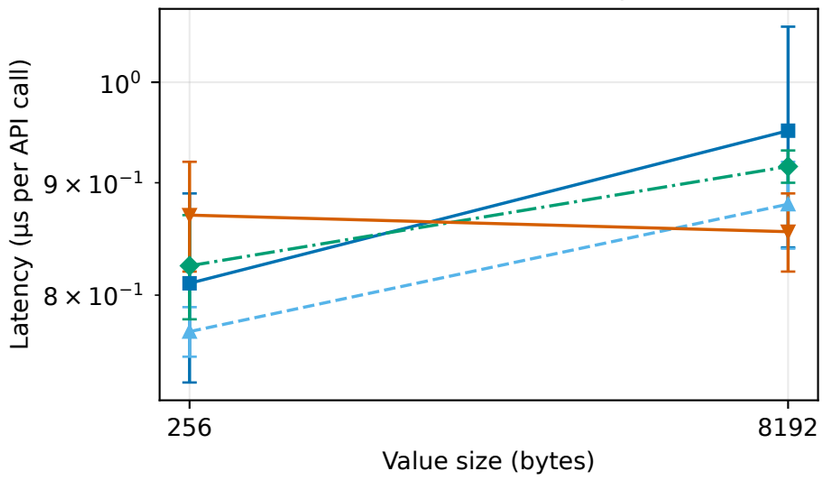

Coordinator cache: API p99

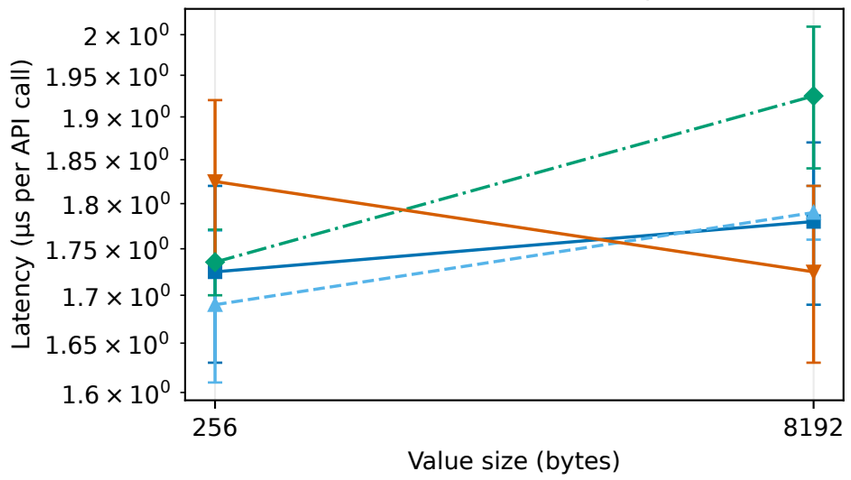

Coordinator cache: successful calls

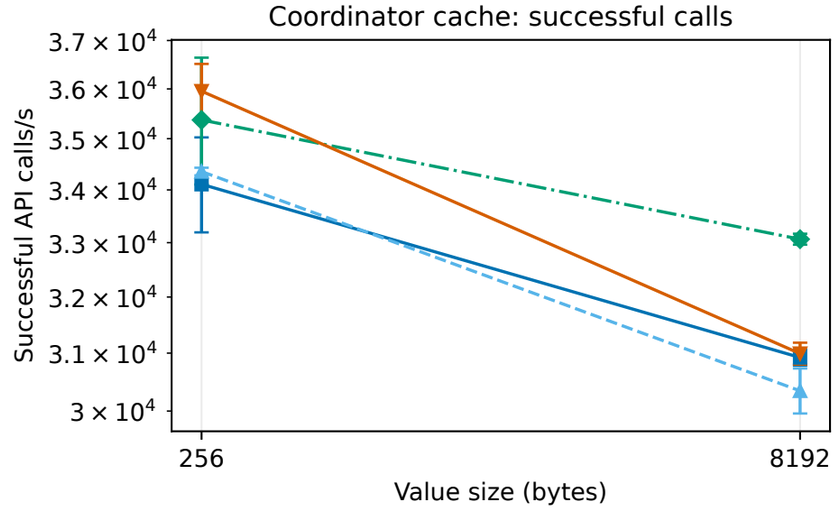

Owner cache: API p50

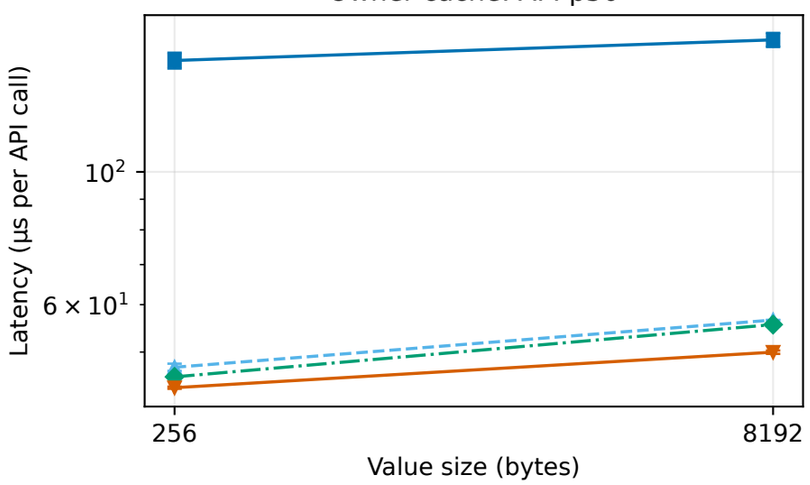

Owner cache: API p99

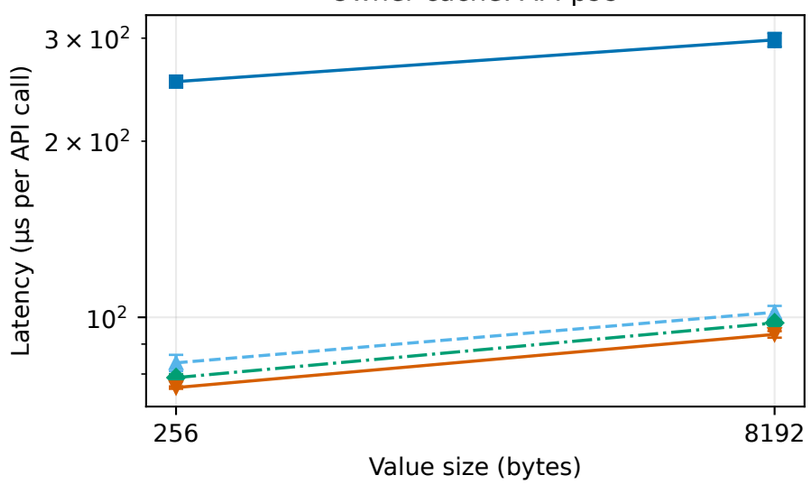

Owner cache: successful calls

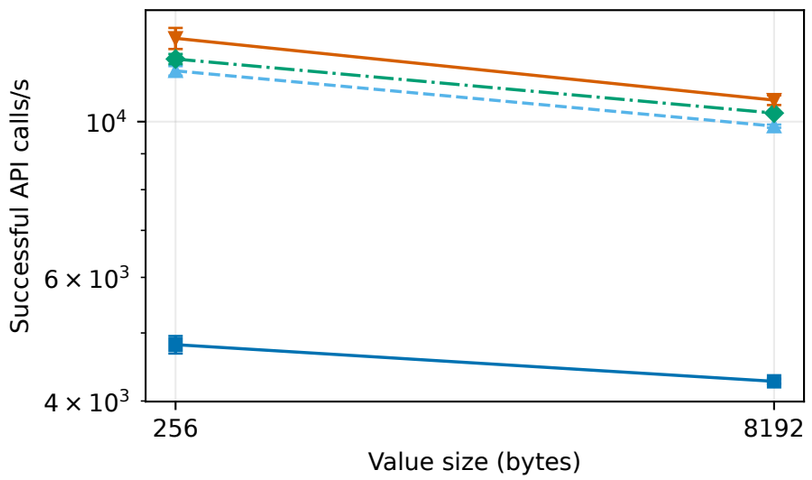

Registered RAM: API p50

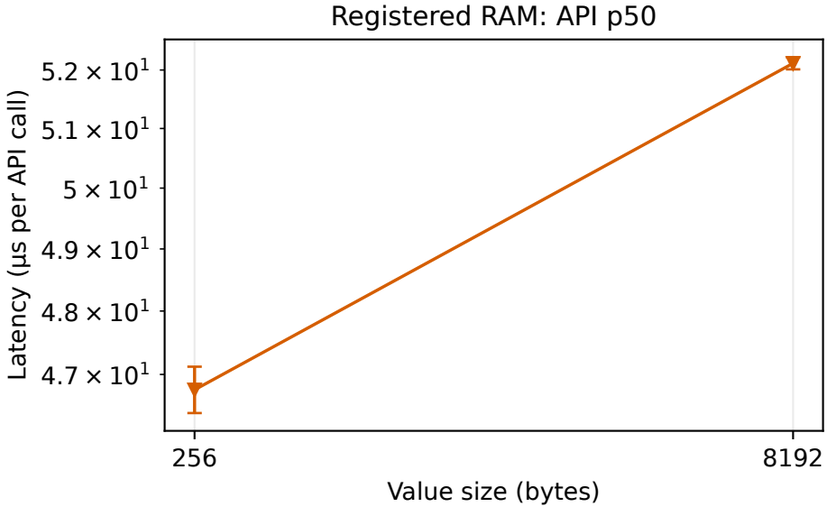

Registered RAM: API p99

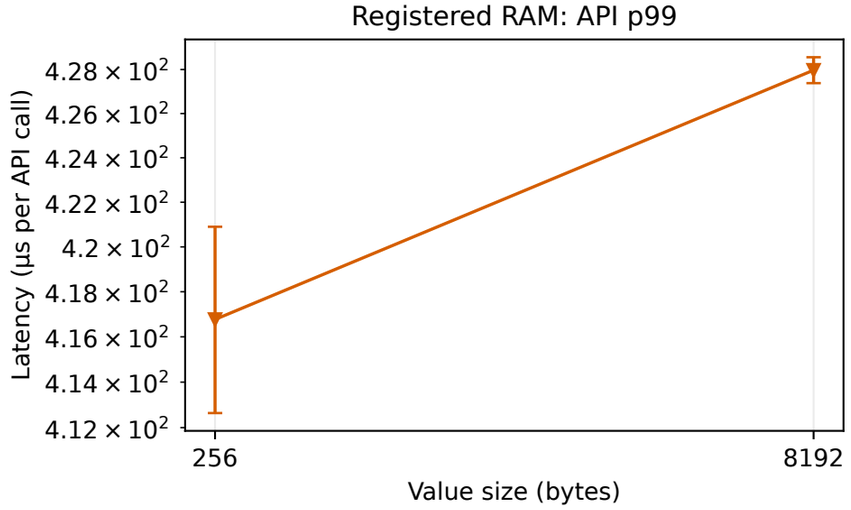

Registered RAM: successful calls

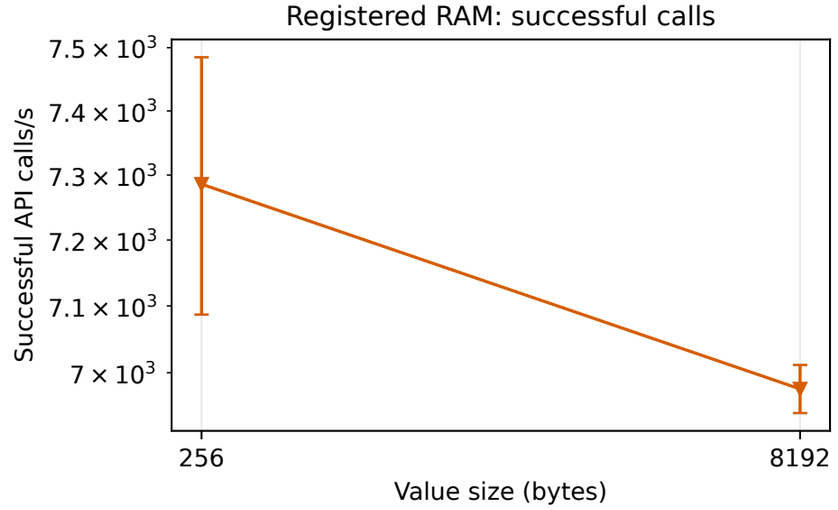

Every row is a distinct c1 cohort; no comparison against primary c16 pure reads. Coordinator and owner cache budgets are 128 MiB where configured; Linux page cache also remains. The owner-cache cohort includes cache misses/WAL fallbacks; it is not isolated owner-cache-hit latency. Registered RAM is TH EFA native only and is mixed RMA plus message fallback, NOT isolated RDMA latency. Cumulative path counters include preload/warmup/qualification and are not timing fractions. Points: median of matched per-repeat values; whiskers: observed run min/max, NOT confidence intervals. No pooled quantiles; few repetitions support no confidence claim. Successful latency is conditional on operation/outcome. RF1 does not equalize consistency, durability, security, or transaction semantics. Primary RF1 only: native uses value_wal_barrier and TH explicit caches are off; other cohorts use their stated settings. Warm OS is not guaranteed cache residency. Tail quantiles are descriptive, including small populations.