Cold-OS-cache protocol — 64 KiB reads, c16
Matched repeats r0, r1

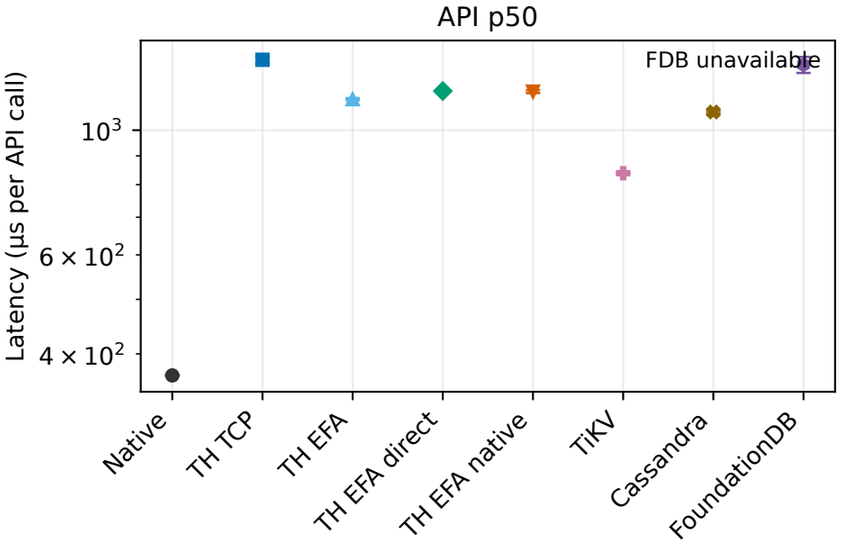

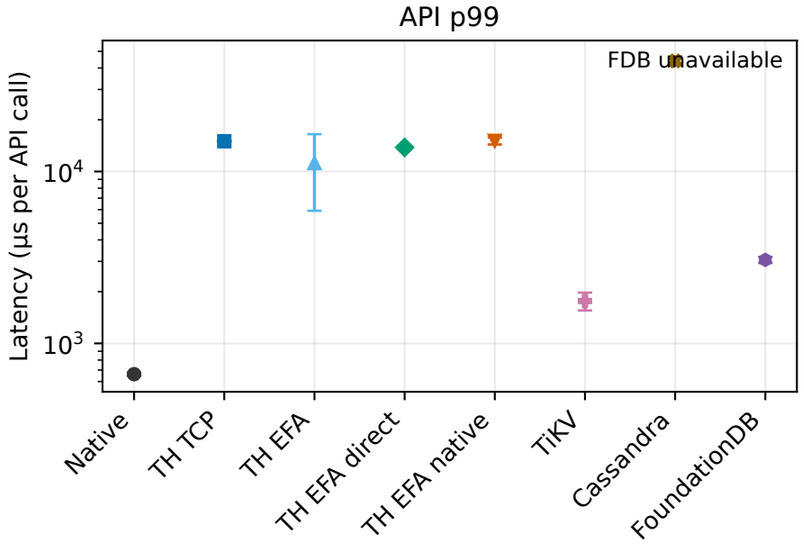

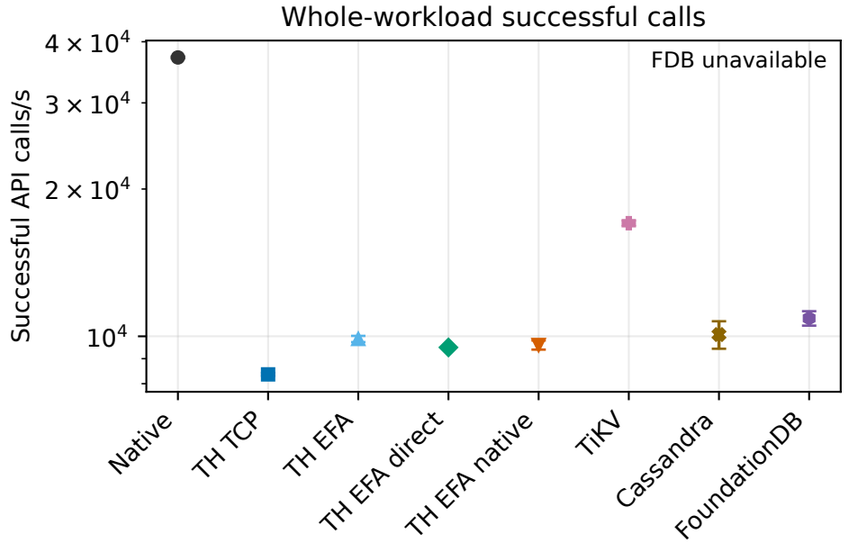

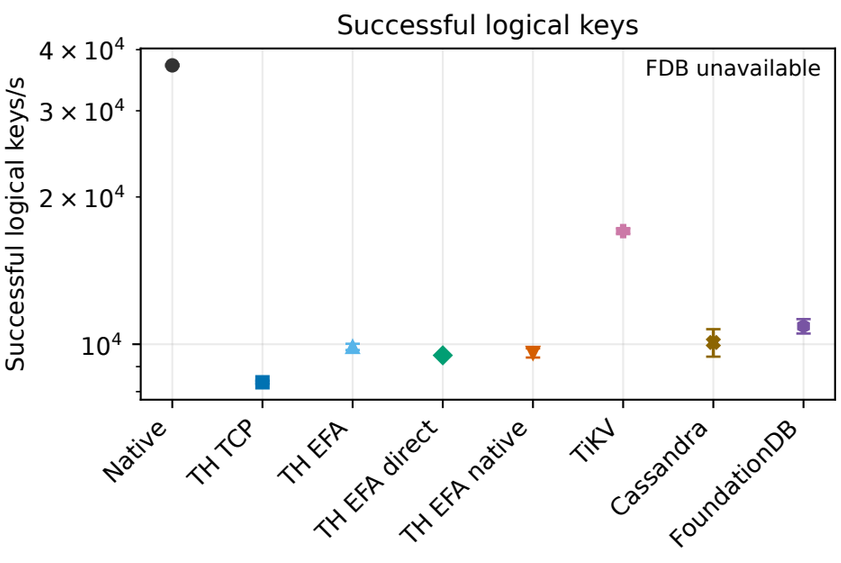

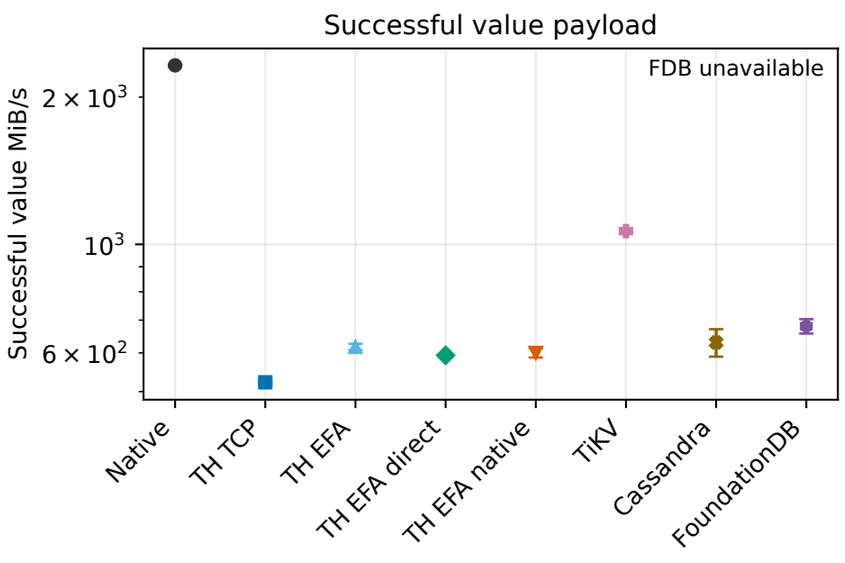

8 GiB value preload, no warmup, cache-drop gate. Cold OS does NOT mean isolated NVMe/device latency: engine and controller caches, network, scheduling and metadata may contribute. Systems need at least two completed cold repeats to enter this separate cohort; all plotted systems share matched repeats and exclusions remain explicit. Missing results are never plotted as zero. Points: median of matched per-repeat values; whiskers: observed run min/max, NOT confidence intervals. No pooled quantiles; few repetitions support no confidence claim. Successful latency is conditional on operation/outcome. RF1 does not equalize consistency, durability, security, or transaction semantics. Primary RF1 only: native uses value_wal_barrier and TH explicit caches are off; other cohorts use their stated settings. Warm OS is not guaranteed cache residency. Tail quantiles are descriptive, including small populations.