Offered-load conditional p99 — 1 KiB, c32
Matched repeats r0, r1

Native
TH TCP
TH EFA
TH EFA direct
TH EFA native
TiKV
Cassandra
FoundationDB

read api p99 / success

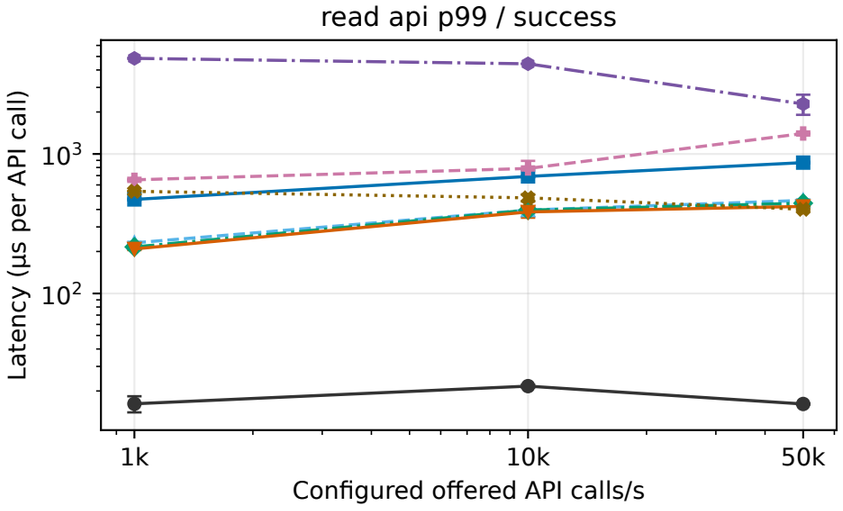

read response p99 / success

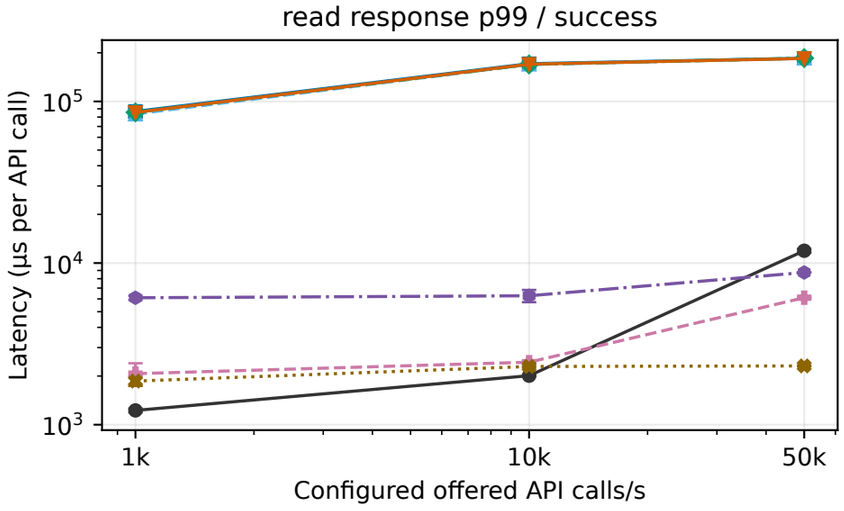

read scheduling p99 / success

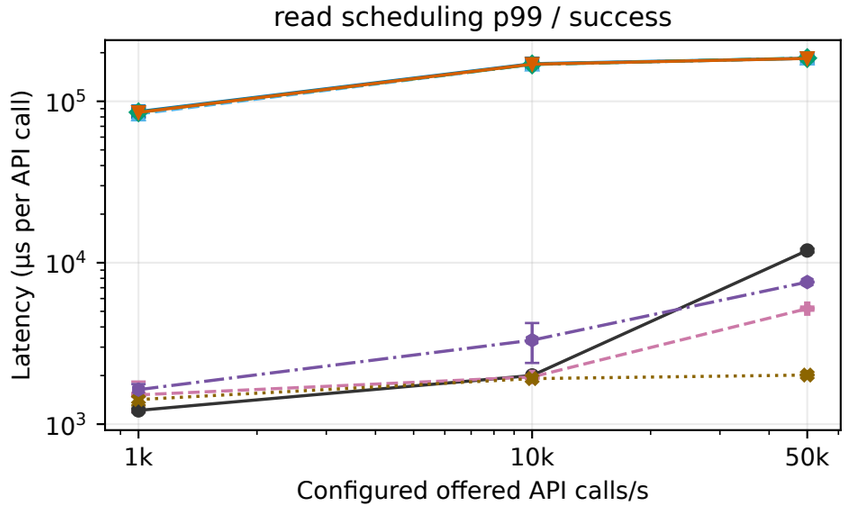

put api p99 / success

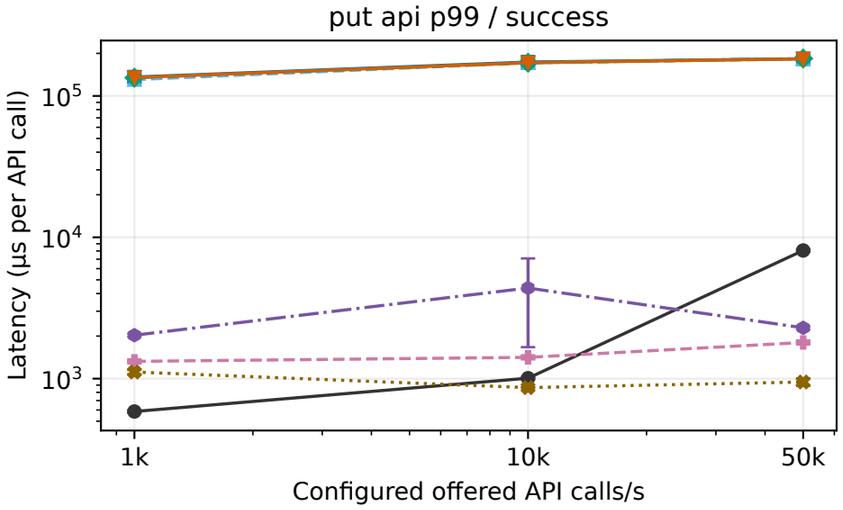

put response p99 / success

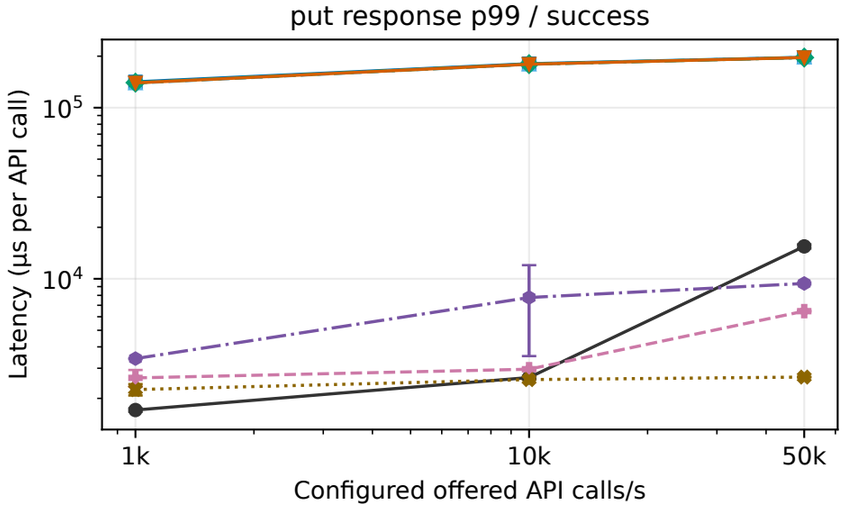

put scheduling p99 / success

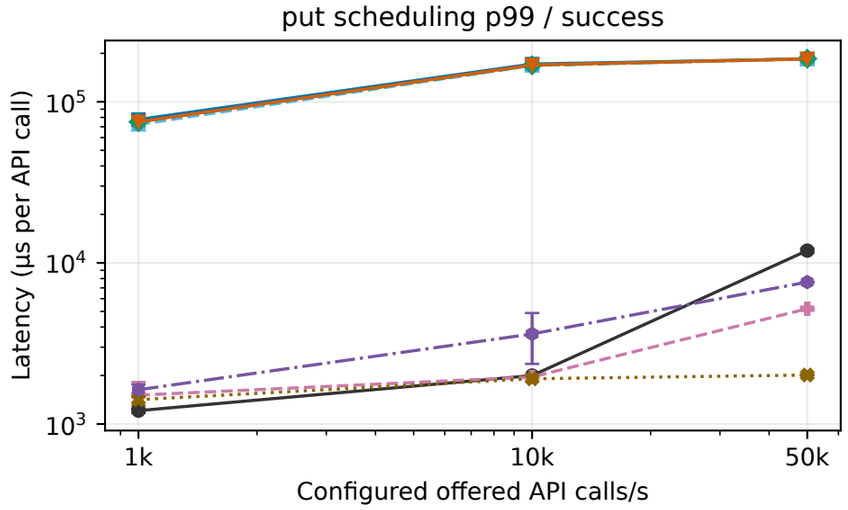

read response p99 / rejected

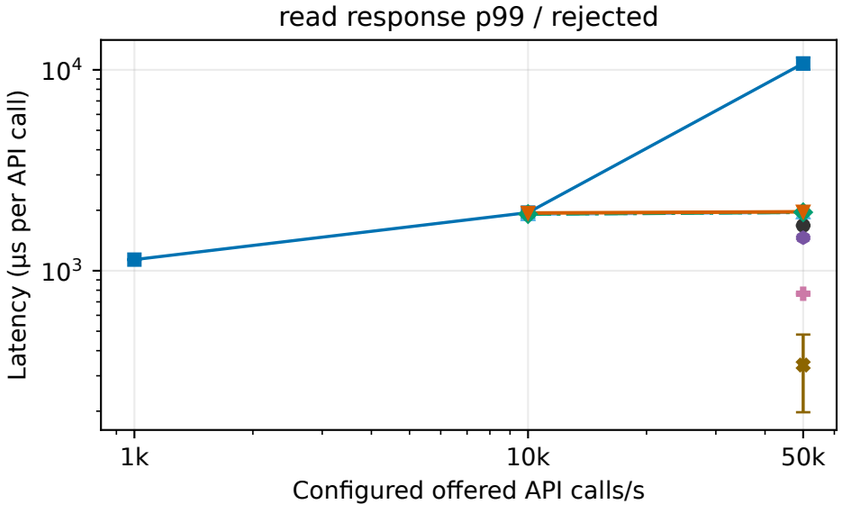

put response p99 / rejected

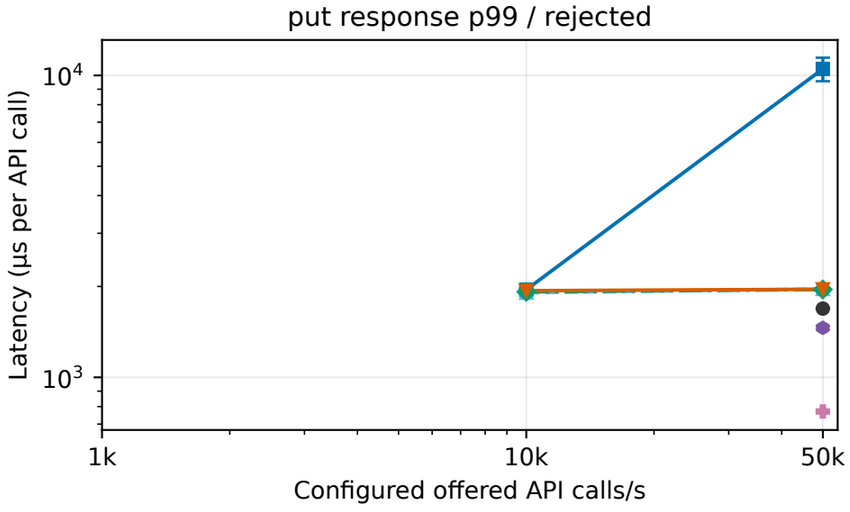

Success and rejected populations are NEVER mixed. API is dispatch-to-completion; response is scheduled-arrival-to-completion; scheduling is arrival-to-dispatch. Rejected arrivals have response only (API/scheduling undefined). Rejected-response points require a timed rejection in every matched repeat; missing populations are gaps, not zeros. Read with offered-load-accounting; fast conditional tails alone do not imply accepting offered load. Points: median of matched per-repeat values; whiskers: observed run min/max, NOT confidence intervals. No pooled quantiles; few repetitions support no confidence claim. Successful latency is conditional on operation/outcome. RF1 does not equalize consistency, durability, security, or transaction semantics. Primary RF1 only: native uses value_wal_barrier and TH explicit caches are off; other cohorts use their stated settings. Warm OS is not guaranteed cache residency. Tail quantiles are descriptive, including small populations.