Cold-OS-cache protocol — 64 KiB reads, c16; 7 available systems
Matched repeats r0, r1

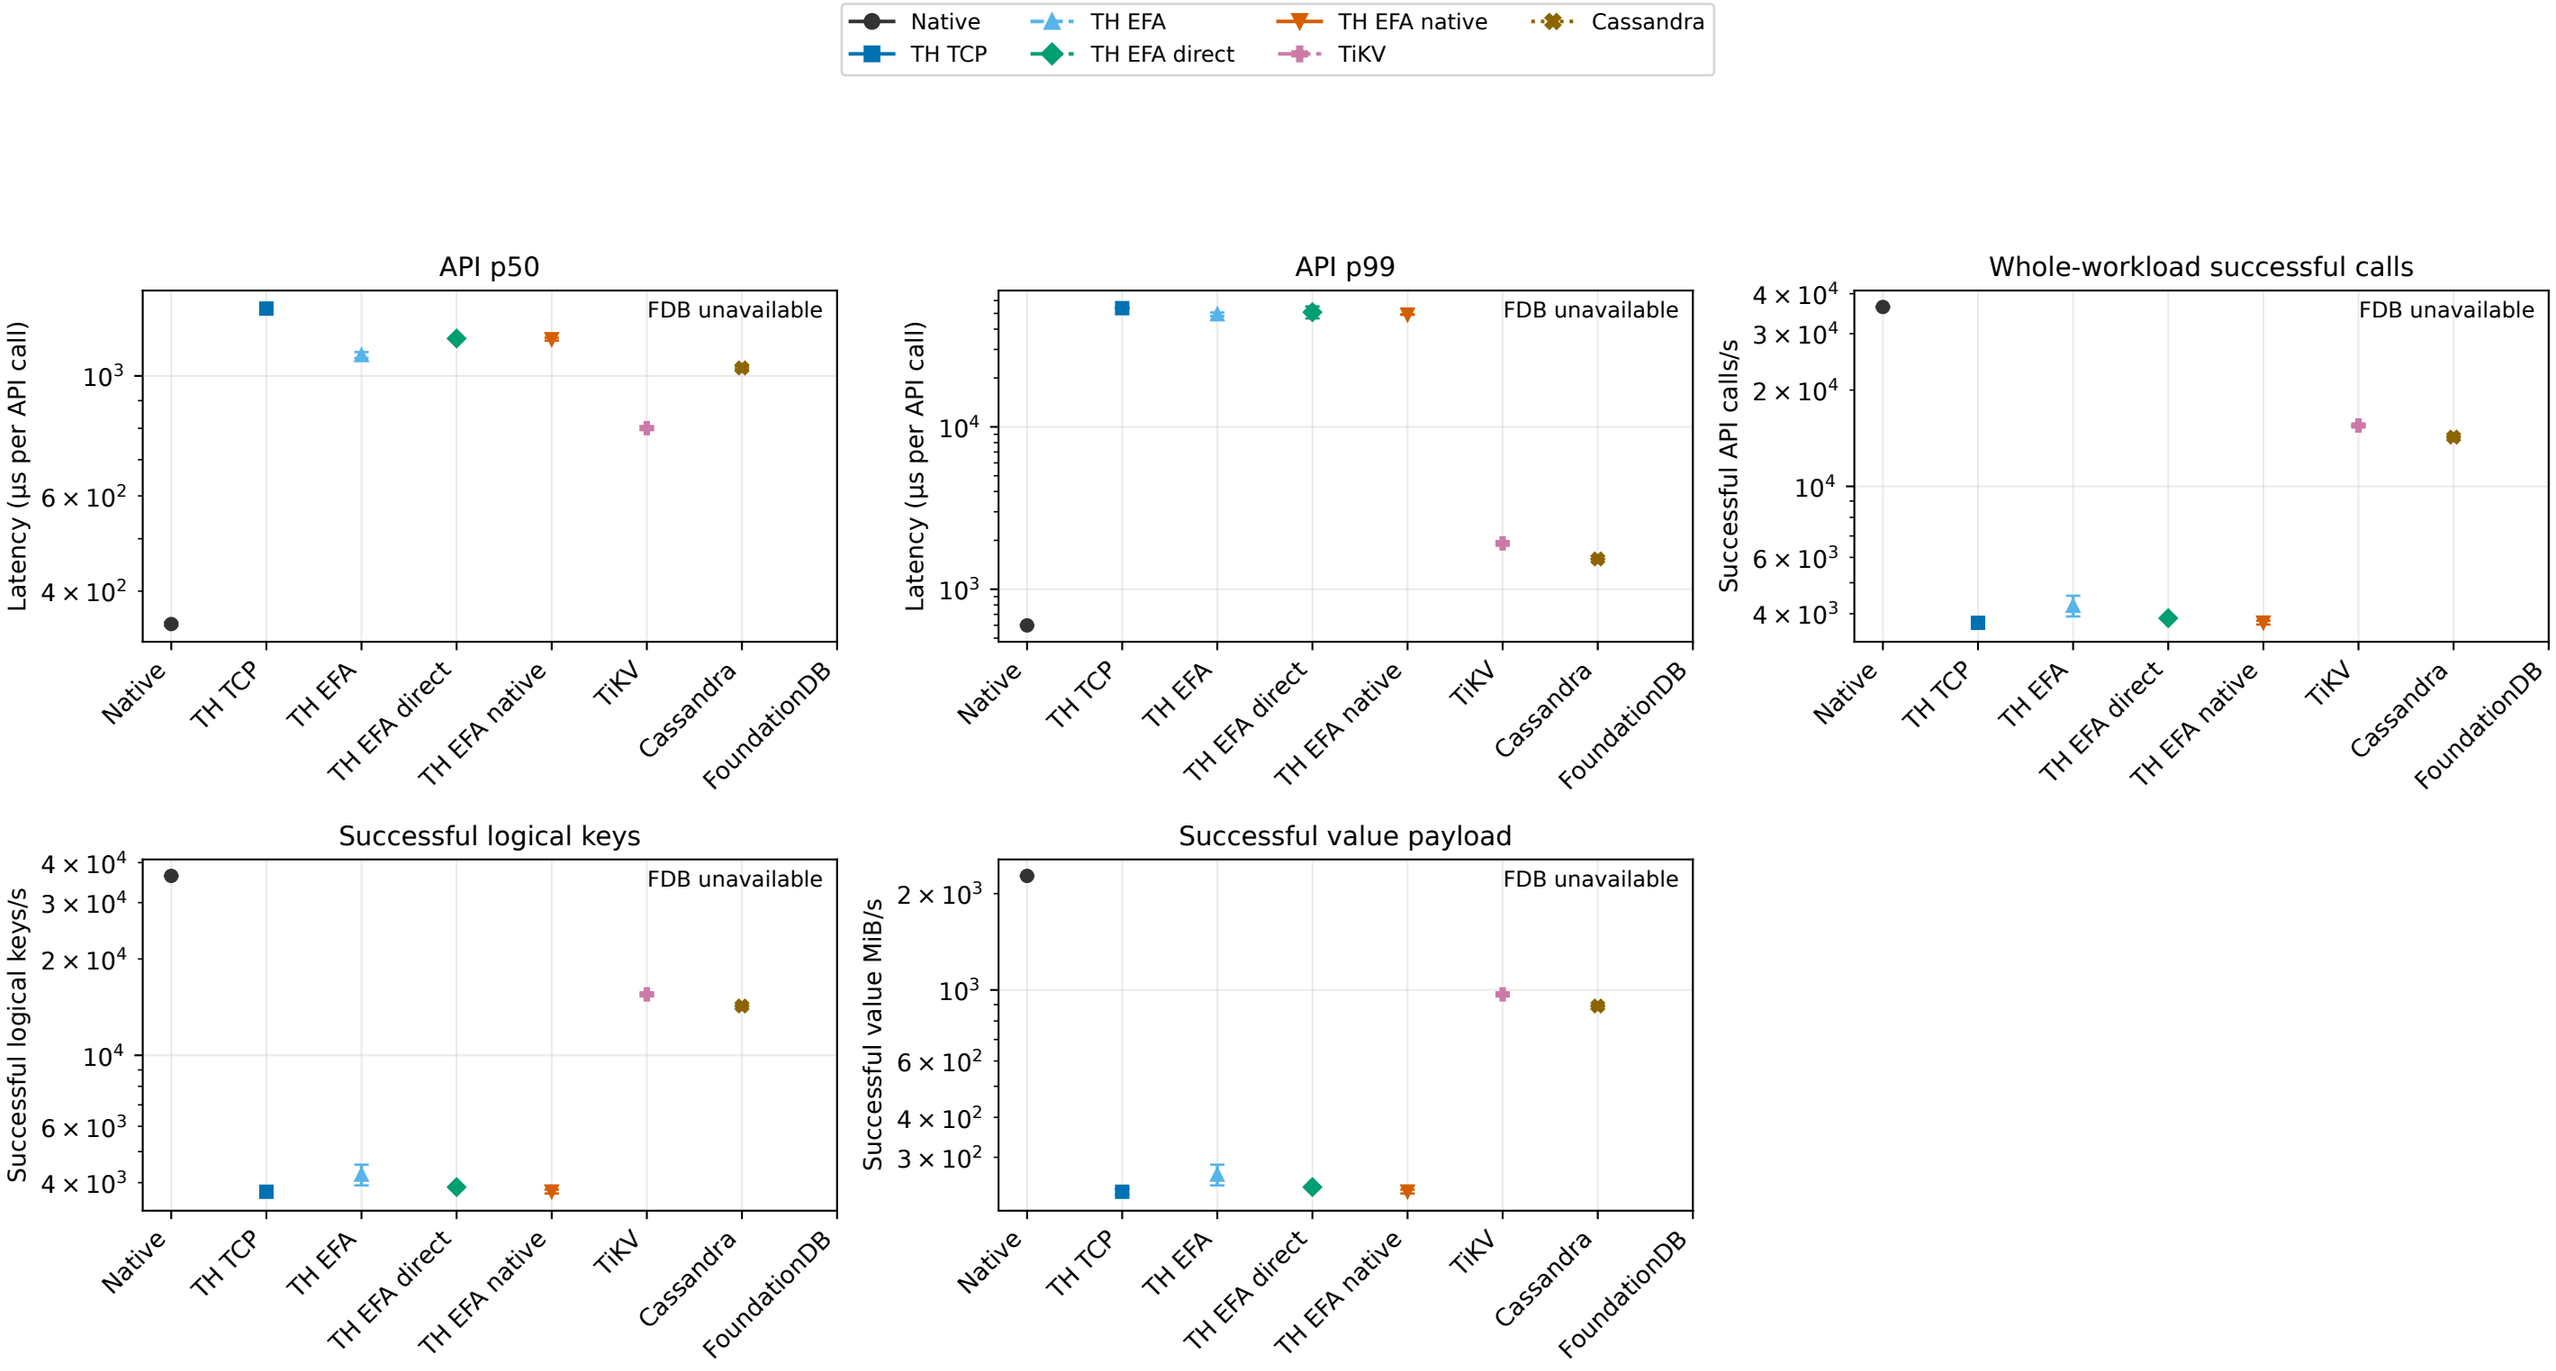

8 GiB value preload, no warmup, cache-drop gate. Cold OS does NOT mean isolated NVMe/device latency: engine and controller caches, network, scheduling and metadata may contribute. FoundationDB unavailable: all three preloads failed before measurement; no latency/goodput zero is invented. Seven-system matched r0/r1 cohort is explicitly separate from primary. Points: median of matched per-repeat values; whiskers: observed run min/max, NOT confidence intervals. No pooled quantiles; N=2 supports no confidence claim. Successful latency is conditional on operation/outcome. RF1 does not equalize consistency, durability, security, or transaction semantics. Primary RF1 only: native uses value_wal_barrier and TH explicit caches are off; other cohorts use their stated settings. Warm OS is not guaranteed cache residency. Tail quantiles are descriptive, including small populations.